



# B3CC: Concurrency

## 04: Threads (2)

---

Ivo Gabe de Wolff

# Last time...

- Processes, threads, threads and threads
- Mutual exclusion
  - Controlling access to shared resources
  - Only one process/thread is allowed in the *critical section* at once

# **Blocking algorithms**

---

# Locks

- A *lock* or *mutex* is a mechanism that enforces limits on access to a resource (mutual exclusion)
  - Conceptually simple!
- Programming languages with support for threads have some form of lock or barrier:
  - Haskell: locking and thread coordination via a mutable data structure called `MVar` (more later...)
  - C/C++: the `std::mutex<T>` class, the POSIX threads library, etc.
  - C#: the `lock` keyword, etc.
  - Rust: the `Mutex<T>` struct
  - etc...

# Locks: historically

- No real hardware support
- e.g. Peterson's algorithm!
- Nowadays: nice hardware instructions!

- Thread A:

```
flag[0] = true;  
turn    = 1;  
  
while (flag[1]  
      && turn == 1)  
/* do nothing */ ;  
  
<critical section>  
  
flag[0] = false;
```

- Thread B:

```
flag[1] = true;  
turn    = 0;  
  
while (flag[0]  
      && turn == 0)  
/* do nothing */ ;  
  
<critical section>  
  
flag[1] = false;
```

# Implementing locks

- The *compare-and-swap* (CAS) ("atomic compare-exchange") operation is an atomic instruction which allows mutual exclusion for any number of threads using a single bit of memory.

```
struct Result {  
    bool success;  
    int original;  
}  
  
Result atomic_compare_exchange(int *variable, int expected, int replacement);
```

- In **hardware**, *atomically* (as a single operation):

1. Compares the contents at a given memory location to the given value
2. If they are the same, writes a new value to that location
3. Returns:
  - whether the new value was written
  - the old value at the memory location

# Implementing locks

- The *spin lock*:

- Use a bit (here 'lock') where 0 represents unlocked and 1 represents locked

```
while (atomic_compare_exchange(&lock, 0, 1).original == 1)
    /* do nothing */ ;

<critical section>

lock = 0;
```

- In Haskell: can use compare-and-swap via casIORef (from atomic-primops)

# The traditional mutex API

- mutex.acquireLock()

```
while (atomic_compare_exchange(&lock, 0, 1).original == 1)  
    /* do nothing */ ;
```

<critical section>

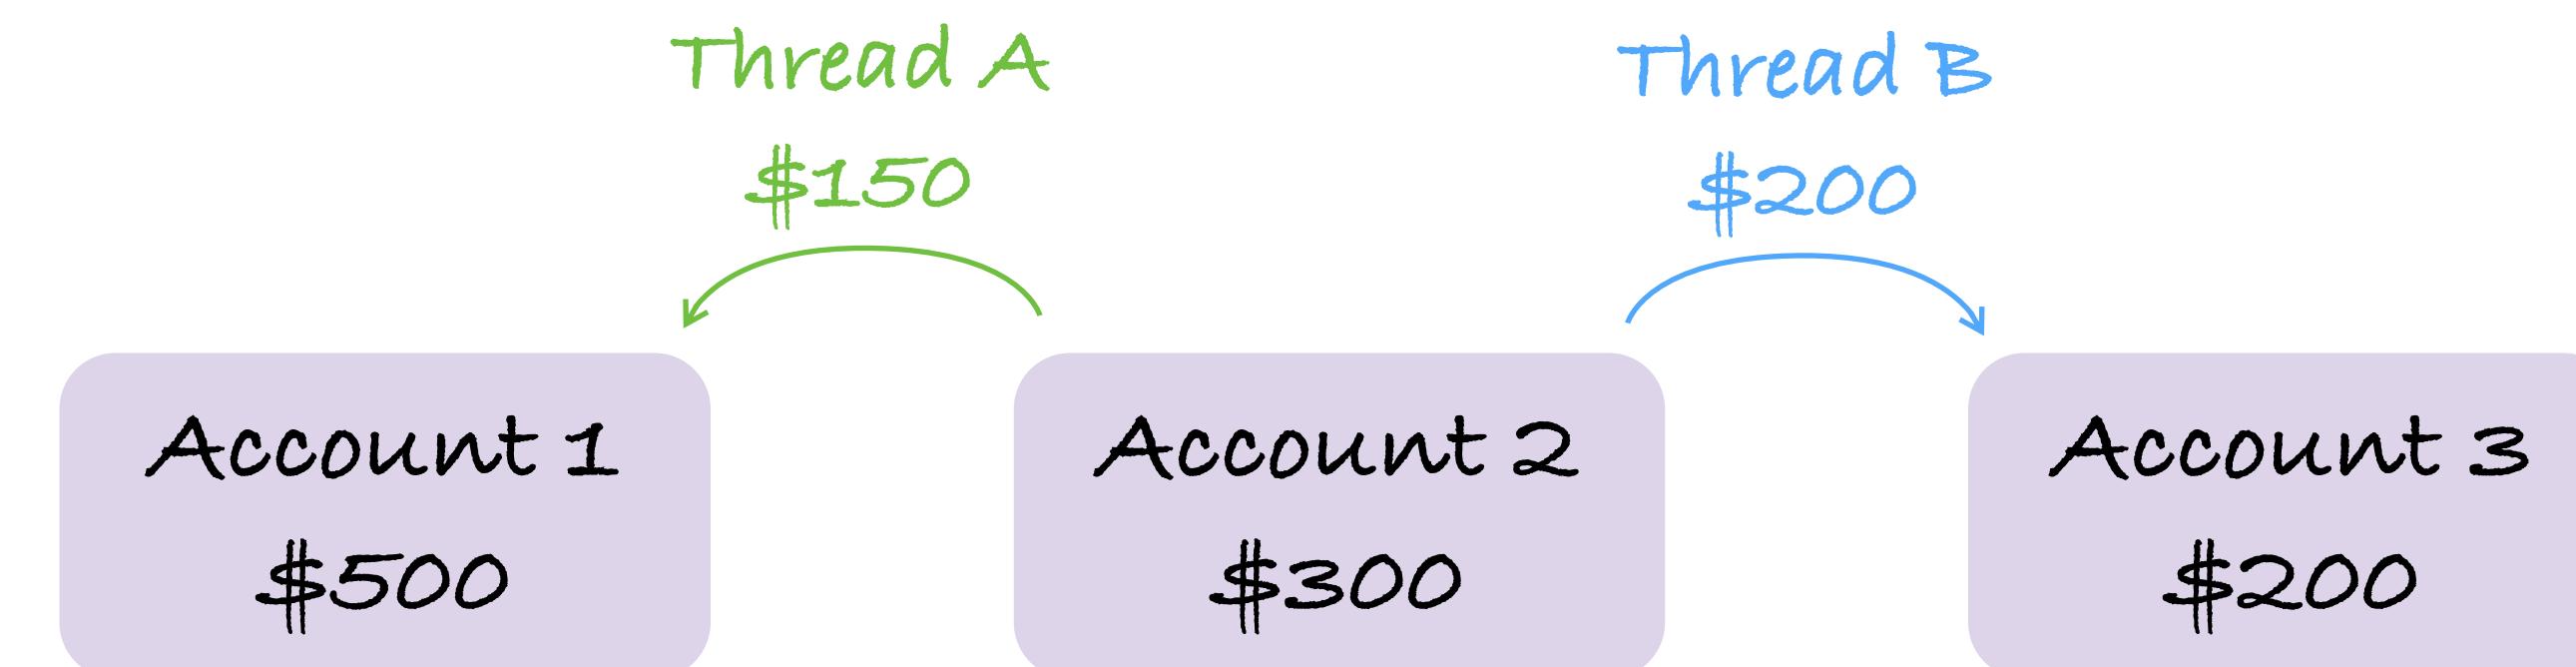
```
lock = 0;
```

- mutex.releaseLock()

- C/C++: pthread\_mutex\_\*, std::mutex<T>; C#: lock; Rust: Mutex<T>; ...

# Example: bank accounts

- Model bank accounts and operations like withdrawing, depositing, and transferring money between accounts
  - It should not be possible to observe a state where, during a transfer, money has been withdrawn from one account but yet to be deposited into the target account



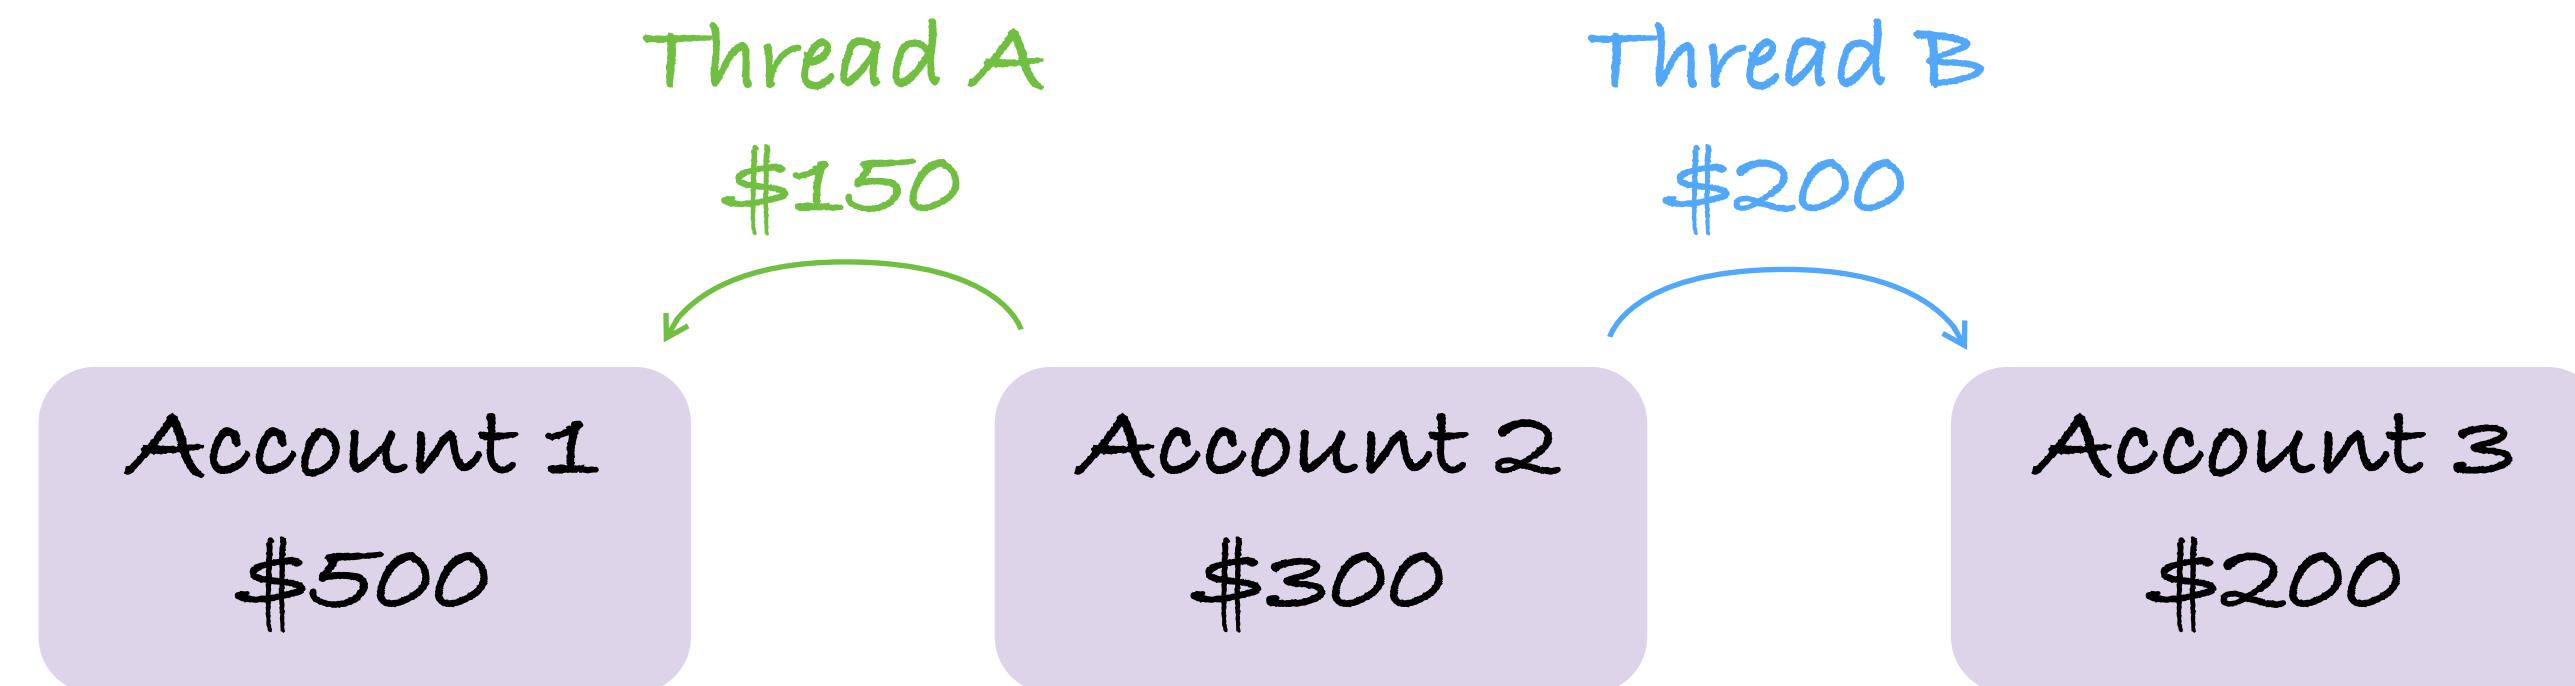
# Attempt #1

- The basic idea:

```
struct Account {  
    int balance;  
};  
  
void deposit(int amount, Account *acc) {  
    int previous = acc->balance;  
    acc->balance = previous + amount;  
}  
  
void withdraw(int amount, Account *acc) {  
    deposit(-amount, acc);  
}  
  
void transfer(int amount, Account *from, Account *to) {  
    withdraw(amount, from);  
    deposit(amount, to);  
}
```

# Example: bank accounts

- Example: bank accounts

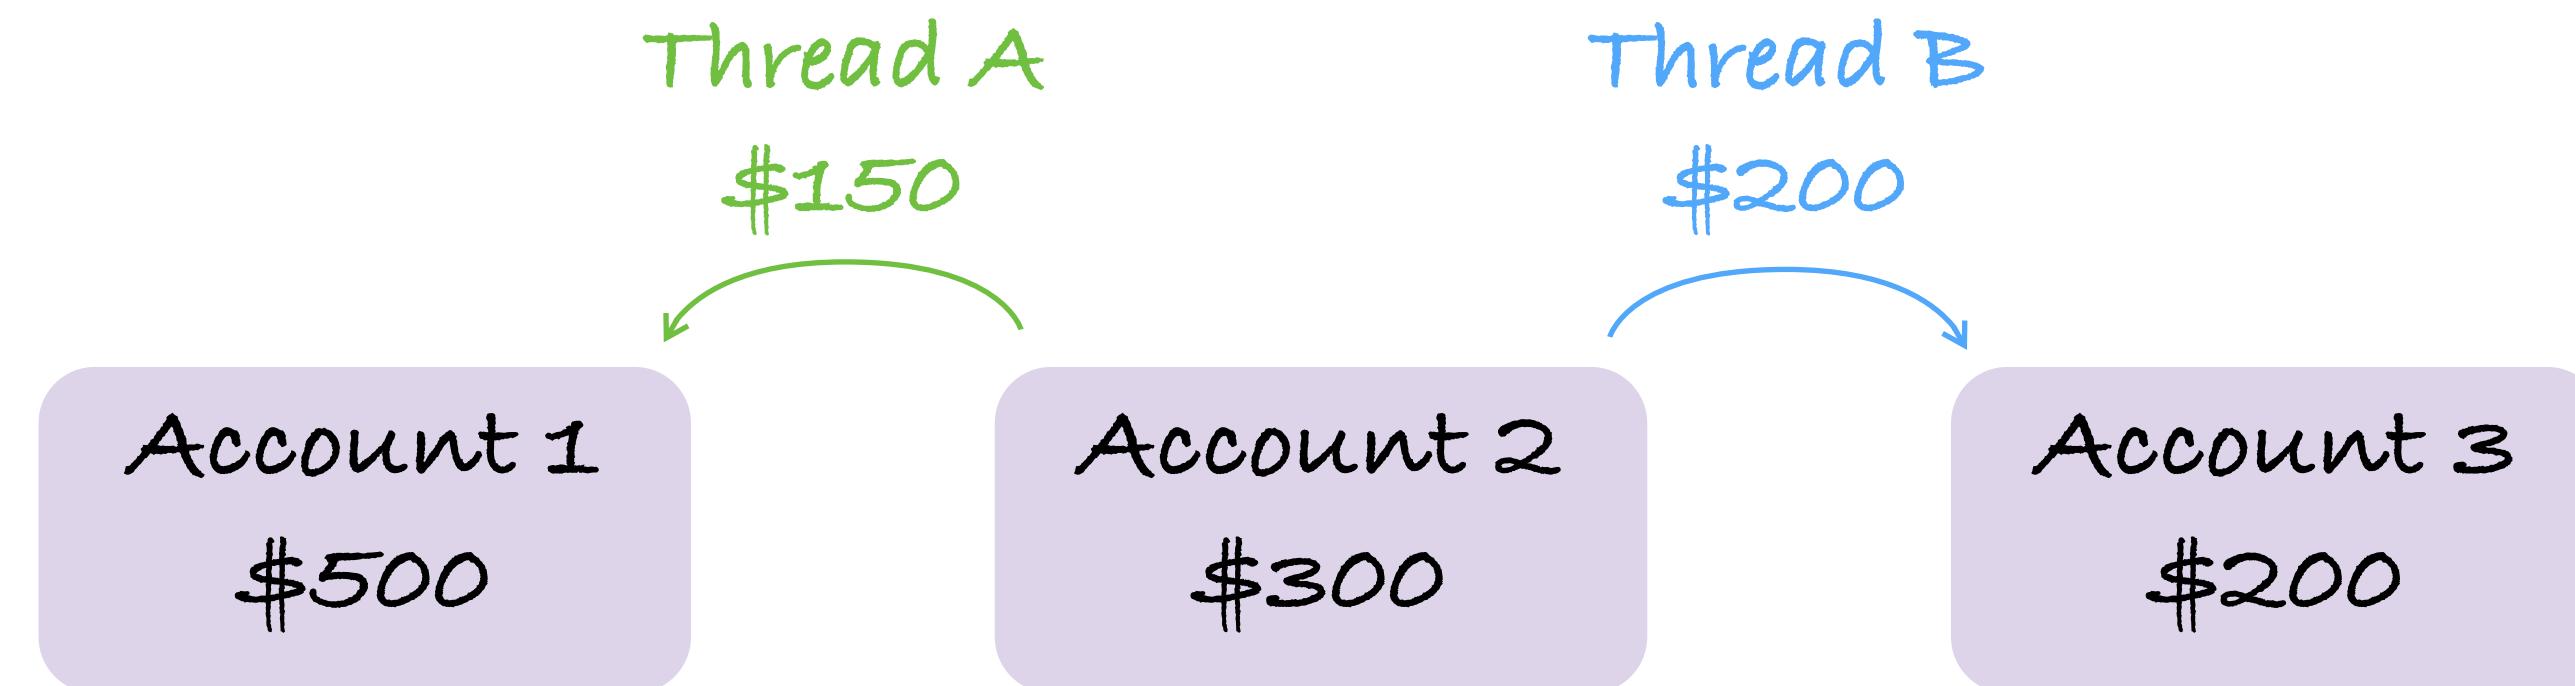


- Thread A:

- Thread B:

# Example: bank accounts

- Example: bank accounts

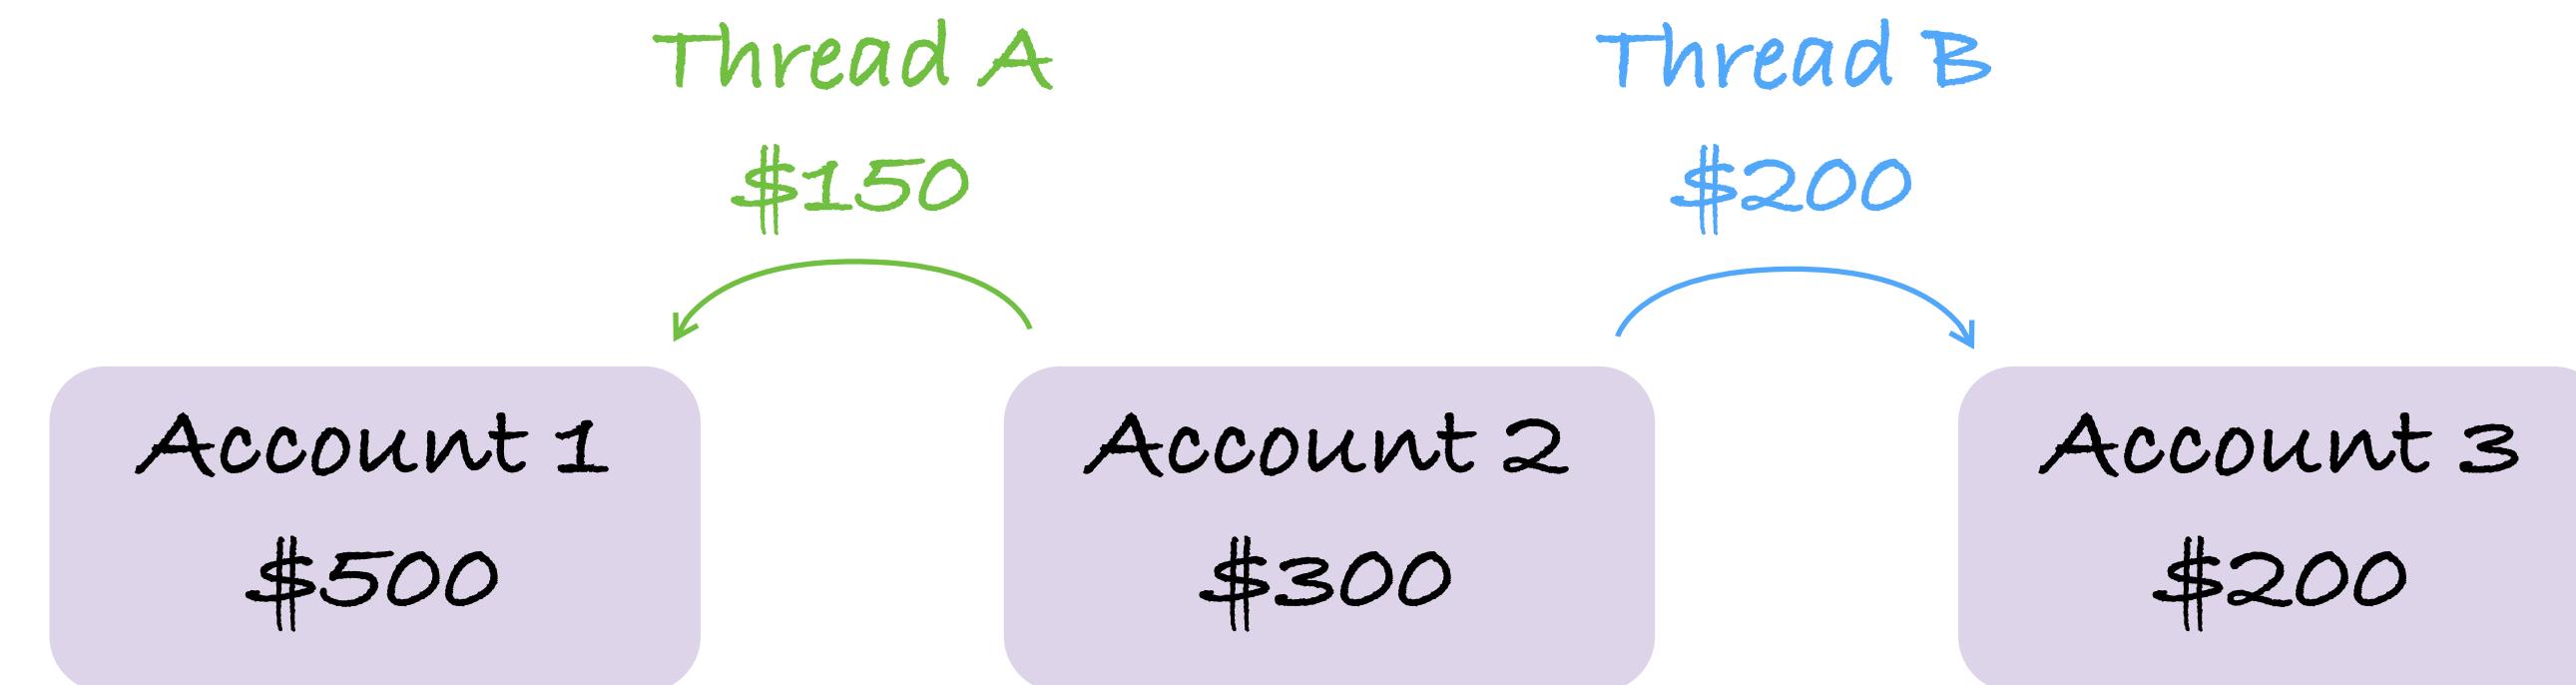


- Thread A:
  - Read balance of account 2: \$300

- Thread B:

# Example: bank accounts

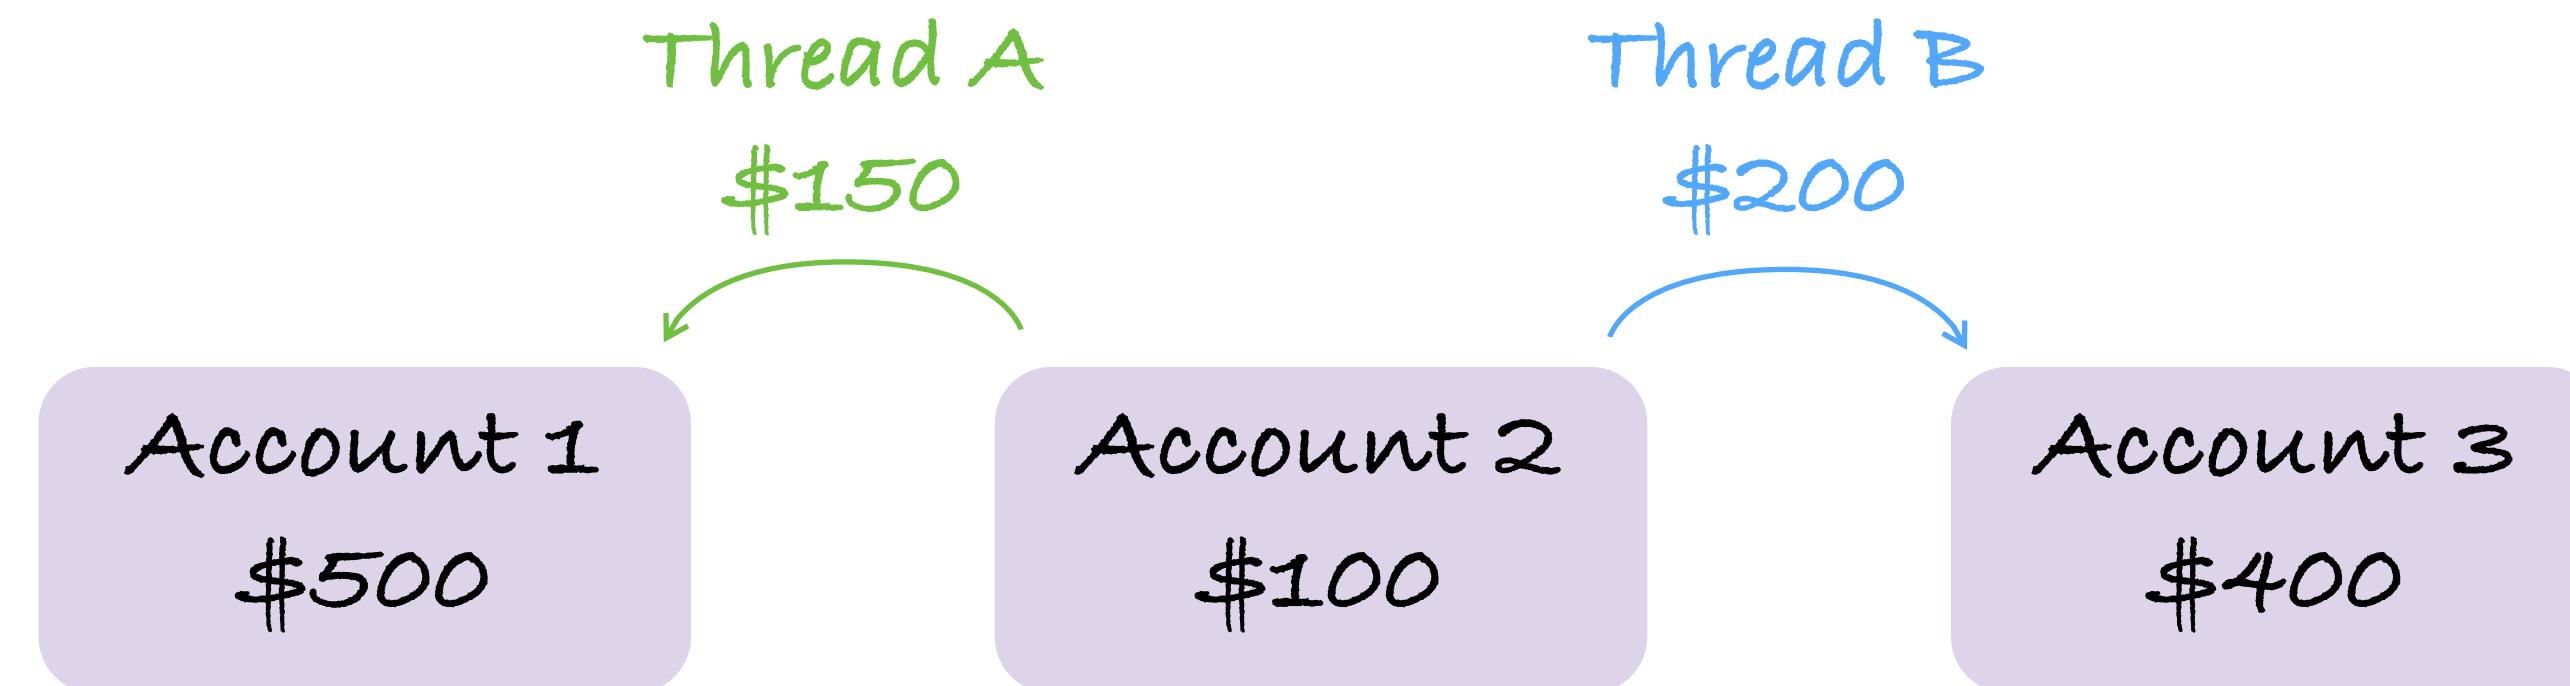
- Example: bank accounts



- Thread A:
  - Read balance of account 2: \$300
- Thread B:
  - Read balance of account 2: \$300

# Example: bank accounts

- Example: bank accounts



- Thread A:
  - Read balance of account 2: \$300
- Thread B:
  - Read balance of account 2: \$300
  - Update balance of account 2

# Example: bank accounts

- Example: bank accounts



- Thread A:
  - Read balance of account 2: \$300
  - Update balance of account 2

- Thread B:
  - Read balance of account 2: \$300
  - Update balance of account 2

## Attempt #2

- Use locks so that updates are atomic:

```
struct Account {  
    int balance;  
    Mutex lock;  
};
```

```
void deposit(int amount, Account *acc) {  
    acc->lock.acquireLock();  
    acc->balance = acc->balance + amount;  
    acc->lock.releaseLock();  
}
```

```
void transfer(int amount, Account *from, Account *to) {  
    withdraw(amount, from);  
    deposit(amount, to);  
}
```

Put balance update  
in a critical section

Let's include a lock this time

Oh no, inconsistent state!

# Attempt #3

- We need to implement transfer differently

```
struct Account {  
    int balance;  
    Mutex lock;  
};
```

```
void transfer(int amount, Account *from, Account *to) {  
    from->lock.acquireLock();  
    to->lock.acquireLock();  
    from->balance = from->balance - amount;  
    to->balance = to->balance + amount;  
    to->lock.releaseLock();  
    from->lock.releaseLock();  
}
```

- Thread A:

- transfer(100, acc1, acc2)

- Thread B:

- transfer(200, acc2, acc1)

# Attempt #4

- Take locks in a fixed (but arbitrary) order; release in the opposite order

```
struct Account {  
    int accountNumber;  
    int balance;  
    Mutex lock;  
};
```

```
void transfer(int amount, Account *from, Account *to) {  
    if (from->accountNumber < to->accountNumber) {  
        from->lock.acquireLock();  
        to->lock.acquireLock();  
        ...  
        to->lock.releaseLock();  
        from->lock.releaseLock();  
    } else {  
        to->lock.acquireLock();  
        from->lock.acquireLock();  
        ...  
        from->lock.releaseLock();  
        to->lock.releaseLock();  
    }  
}
```

# Extending the example

- What happens if we want to...
  - Block (wait) until the 'from' account has sufficient funds?
  - Withdraw from a second account if the first does not have sufficient funds?
    - Suppose I hold locks #3 and #5...
    - And now need to acquire lock #2, or #4, or...

# Advantages and disadvantages

- Difficulties / problems (among others):
  - Taking locks in the wrong order
  - Too few locks (*lock contention decreases the amount of available concurrency*)
  - Too many locks (increases overhead and subtle lock dependencies that can increase the chance of *deadlock*)
  - Error recovery (out of scope for this course)
  - No modular programming! (transfer; lock order)
- Advantages:
  - Easy critical sections if you have a single lock
  - Mutual exclusion!

*User space*

## **Threads in Haskell**

---

# Threads

- The fundamental action in concurrency: create a new thread of control

```
forkIO :: IO () -> IO ThreadId
```

- Takes a computation of type `IO ()` as its argument
- This `IO` action executes in a new thread concurrently with other threads
- No specified order in which threads execute
- Haskell user space threads are very cheap: ~1.5 KB / thread, easily run thousands of threads

# Example

- Interleaving of two threads

```
import Control.Concurrent
import Control.Monad

main :: IO ()
main = do
    let n = 100

    forkIO $ replicateM_ n (putChar 'A')
    forkIO $ replicateM_ n (putChar 'x')

    putStrLn "done"
```

- Interleaving of two threads

- The program exits when main returns, even if there are other threads still running!
  - How to check whether the child thread has completed?
- The term n :: Int is shared between both threads (captured); this is safe because it is *immutable*

# **Sharing state**

---

# Sharing state

- `I0Ref`: mutable reference to some value
  - i.e. a regular variable in C#
  - Compare-and-swap behaviour using `casI0Ref`
  - Not designed for concurrency: need to protect critical sections yourself
- `MVar`: *synchronised* mutable references
  - Like `I0Refs`, but with a *lock* attached for safe concurrent access
  - ...plus some very useful semantics

# IORefs

- In most languages variables are mutable by default
- In Haskell, mutable variables must be handled explicitly
  - Notice that whether a variable is mutable is now reflected in its type!

```
import Data.IORef

newIORef    :: a -> IO (IORef a)
readIORef   :: IORef a -> IO a
writeIORef  :: IORef a -> a -> IO ()
```

# IORefs: Example

- Shared state concurrency using IORef

```
import Control.Concurrent
import Data.IORef

main :: IO ()
main = do
    ref <- newIORef 0

    forkIO $ writeIORef ref 1
    forkIO $ writeIORef ref 2

    result <- readIORef ref -- `_(ツ)_/
    print result
```

# MVars

- Synchronising variables for communication between concurrent threads
  - An MVar is a box that is either *empty* or *full*
  - `takeMVar` removes the value from the box; blocks if it is currently empty
  - `putMVar` puts a value in the box; blocks if it is currently full
  - `readMVar` reads the current value without removing it (and blocks if empty)

```
import Control.Concurrent

newMVar      :: a -> IO (MVar a)
newEmptyMVar :: IO (MVar a)
takeMVar     :: MVar a -> IO a
putMVar      :: MVar a -> a -> IO ()
readMVar     :: MVar a -> IO a
```

# MVars

- Synchronising variables for communication between concurrent threads

```
import Control.Concurrent

main :: IO ()
main = do
    m <- newEmptyMVar
    forkIO $ do
        putMVar m "hello"
        putMVar m "world"

    x <- takeMVar m
    putStrLn x
    putStrLn ","
    y <- takeMVar m
    putStrLn y
```

# MVars

- The runtime system can (sometimes) detect when a group of threads are deadlocked
  - Only a conservative approximation to the future behaviour of the program
  - Can be useful for debugging (but don't rely on it)

```
main :: IO ()  
main = do  
    m <- newEmptyMVar  
    takeMVar m
```

```
$ ./Test  
Test: thread blocked indefinitely in an MVar operation
```

# MVars

- If multiple threads are blocked in `takeMVar` or `putMVar`, a *single thread* is woken up in FIFO order: fairness
- If `readMVar` blocks, it will receive the *next* put value
- Other useful operations
  - `withMVar` can be used to protect critical sections (read the docs!)

```
takeMVar      :: MVar a -> IO a
putMVar       :: MVar a -> a -> IO ()
readMVar      :: MVar a -> IO a

withMVar      :: MVar a -> (a -> IO b) -> IO b
```

# An MVar is...

- A lock
  - `MVar ()` behaves as a lock: full is unlocked, empty is locked
  - Can be used as a mutex to protect some shared state or critical section
- A one-place channel
  - For passing messages between threads
  - An asynchronous channel with a buffer size of one
- A container for shared mutable data
- A *building block* for constructing larger concurrent data structures

# **MVars as a building block (I)**

---

Asynchronous computations

# Asynchronous computations

- The goal:
  - Want a way to run computations asynchronously and wait for their results
  - Cancel running computations
  - Basic interface:

```
data Async a

runAsync :: IO a -> IO (Async a)
wait      :: Async a -> IO a
poll      :: Async a -> IO (Maybe a)
cancel    :: Async a -> IO ()
```

# runAsync

- Perform an action *asynchronously*, and later wait for the results

```
data Async a = Async ThreadId (MVar a)

runAsync :: IO a -> IO (Async a)
runAsync action = do
    var <- newEmptyMVar
    tid <- forkIO $ do
        res <- action
        putMVar var res

    return (Async tid var)
```

# wait, cancel and poll

- Wait for the computation to complete or cancel it

```
wait :: Async a -> IO a
wait (Async _ var) = readMVar var
```

```
cancel :: Async a -> IO ()
cancel (Async tid _) = killThread tid
```

```
poll :: Async a -> IO (Maybe a)
poll (Async _ var) = tryReadMVar var
```

# Cancelling a thread

- `killThread :: ThreadId -> IO ()`
- May lead to subtle bugs
  - What if the killed thread had a lock and was in the critical section?
  - We won't use `killThread` in this course
- Alternative: cancellation token
  - For instance, `IORef Bool`
  - Write `True` if you want to stop a task
  - Check regularly (at safe points) whether threads should stop

# IBAN

- You can now implement count and list in P1
- Search mode after Thursday's lecture
- There may be questions about the practical assignments on the exams

# Next time

- Non-blocking algorithms
- IORefs and MVars as building blocks