

# B3CC: Concurrency

## 06: Software Transactional Memory (I)

---


Tom Smeding

# **Recap**

---

# Why?

- Concurrency control required for safe access to shared state between threads
  - Examples we've seen previously:



# Attempt #4

- Take locks in a fixed (but arbitrary) order; release in the opposite order

```
struct Account {  
    int balance;  
    Mutex lock;  
};
```

```
void transfer(int amount, Account *from, Account *to) {  
    if (from->accountNumber < to->accountNumber) {  
        from->lock.acquireLock();  
        to->lock.acquireLock();  
        ...  
        to->lock.releaseLock();  
        from->lock.releaseLock();  
    } else {  
        to->lock.acquireLock();  
        from->lock.acquireLock();  
        ...  
        from->lock.releaseLock();  
        to->lock.releaseLock();  
    }  
}
```

LAST WEEK

# Why?

- Concurrency control
  - *Mutual exclusion*: critical resources => critical section
    - Only one process allowed in the critical section at once
  - *Deadlock*
  - *Starvation*

# Review

- What are the requirements for *implementing* mutual exclusion?
- What are the requirements for *using* critical sections?

# Review


- Using critical sections
  - Threads should stay in the critical section for as little time as possible
  - What is the consequence of taking locks for too long?

```
countMode :: MVar Int -> [Int] -> IO ()  
countMode var accounts =  
    sum [ 1 | a <- accounts, mtest m a ]
```



# Dining philosophers

- Canonical example of synchronisation issues and how to resolve them
  - Philosophers alternatively think and eat
  - Require both forks to start eating
  - Each fork is held by one philosopher at a time



# **Atomic blocks**

---

# An alternative

- The idea:
  - Garbage collectors allow us to program without `malloc()` and `free()`
    - Can we do the same for locks?
    - What would that look like?
  - Modular concurrency!
  - Locks are pessimistic; let's be optimistic instead!

# Software transactional memory

- A [programming languages/software-based] technique for implementing *atomic blocks*
  - Atomicity: effects become visible to other threads all at once
  - Isolation: cannot see the effects of other threads
  - Use a different type (STM) to wrap operations whose effects *can be undone if necessary* (more on this later)

```
import Control.Concurrent.STM

data STM a           -- abstract
instance Monad STM  -- among other things

atomically :: STM a -> IO a
```

# Software transactional memory

- Sharing state
  - Instead of `IORef`, we use `TVar` as a *transactional variable*
  - Basic interface:

```
import Control.Concurrent.STM.TVar

newTVar    :: a -> STM (TVar a)
readTVar   :: TVar a -> STM a
writeTVar :: TVar a -> a -> STM ()
```

# Revisiting accounts

- STM actions are composed together in the same way as IO actions

```
type Account = TVar Int

deposit :: Int -> Account -> STM ()
deposit amount account = do
    balance <- readTVar account
    writeTVar account (balance + amount)

withdraw :: Int -> Account -> STM ()
withdraw amount = deposit (-amount)

transfer :: Int -> Account -> Account -> IO ()
transfer amount from to =
    atomically $ do
        withdraw amount from
        deposit amount to
```

```
void transfer(int amount, Account *from, Account
if (from->accountNumber < to->accountNumber) {
    from->lock.acquireLock();
    to->lock.acquireLock();
    ...
    to->lock.releaseLock();
    from->lock.releaseLock();
} else {
    to->lock.acquireLock();
    from->lock.acquireLock();
    ...
    from->lock.releaseLock();
    to->lock.releaseLock();
}
}
```

# STM

- Types are used to isolate transactional actions from arbitrary IO actions
  - To get from STM to IO we have to execute the entire action atomically
  - Can't mix monads!

```
bad :: Int -> Account -> ?? ()  
bad amount account = do  
    putStrLn "withdrawning!"  
    withdraw amount account
```

```
good :: Int -> Account -> IO ()  
good amount account = do  
    putStrLn "withdrawning!" -- :: IO ()  
    atomically $ withdraw amount account -- :: IO ()
```

# Implementing transactional memory

- How to implement atomically
  - Single global lock?
  - Instead: optimistic execution, without taking any locks
- At the start of the atomic block begin a thread local *transaction log*
  - Each writeTVar records the address and the new value to the log
  - Each readTVar searches the log and
    - Takes the value of an earlier readTVar / writeTVar; or
    - Reads the TVar and records the value into the log

# Implementing transactional memory

- At the end of the atomic block the transaction log must be *validated*
  - Checks each `readTVar` in the log matches the current value
  - If successful all `writeTVar` recorded in the log are *committed* to the real `TVars`
  - The validate and commit steps together must be truly atomic

# Implementing transactional memory

- What if validation fails?
  - The operation executed with an inconsistent view of memory
  - Re-execute the transaction with a new transaction log
    - Since none of the writes are committed to memory, this is safe to do
    - It is critical that the atomic block contains no actions other than reads and writes to TVars

```
atomically $ do
    x <- readTVar xv
    y <- readTVar yv
    if x > y
        then system "rm -rf /"      -- :: IO () side effects!
    else return ()
```

# Summary (so far)

- STM gives us:
  - Atomic transactions for shared memory
  - Encapsulation of concurrent code
  - Help avoid common locking problems
- Locks are pessimistic, STM is optimistic
- But...
  - Just like garbage collection, is no silver bullet: blocking; starvation & contention

# **Blocking & Choice**

---

# Software transactional memory

- Sharing state
  - Instead of MVar we have an equivalent TMVar
  - A variable is either *full* or *empty*: threads wait for the appropriate state
  - Basic interface:

```
import Control.Concurrent.STM.TMVar

newTMVar      :: a -> STM (TMVar a)
newEmptyTMVar :: STM (TMVar a)
takeTMVar     :: TMVar a -> STM a
readTMVar     :: TMVar a -> STM a
putTMVar      :: TMVar a -> a -> STM ()
```

# Blocking

- Wait for some condition to be true or a resource to become available
  - Abandon the current transaction and begin again
  - Only when the inputs change, to avoid busy waiting (how?)

```
retry :: STM a
```

# Accounts, revisited

- Suppose we want to block if the account will be overdrawn
  - Because the transaction read account on the way to retry, the thread can wait until this variable changes before trying again

```
type Account = TVar Int

withdraw :: Int -> Account -> STM ()
withdraw amount account = do
    balance <- readTVar account
    if amount > 0 && amount > balance
        then retry
        else writeTVar account (balance - amount)
```

# Example: TMVar

- Transactional equivalent of MVar
  - Shared variable which is either empty or full
  - Easy to implement in terms of TVar!

```
newtype TMVar a = TMVar (TVar (Maybe a))
```

```
newEmptyTMVar :: STM (TMVar a)
takeTMVar      :: TMVar a -> STM a
putTMVar       :: TMVar a -> a -> STM ()
```

```
newEmptyTMVar :: STM (TMVar a)
newEmptyTMVar = do
    t <- newTVar Nothing
    return (TMVar t)
```

# TMVar

- Block if the desired variable is empty, and return the contents when it is full

```
takeTMVar :: TMVar a -> STM a
takeTMVar (TMVar t) = do
    m <- readTVar t
    case m of
        Nothing -> retry
        Just a -> do
            writeTVar t Nothing
            return a
```

```
newtype TMVar a = TMVar (TVar (Maybe a))
```

# TMVar

- Block when the variable is full, update the contents when it is empty

```
putTMVar :: TMVar a -> a -> STM a
putTMVar (TMVar t) a = do
    m <- readTVar t
    case m of
        Nothing -> writeTVar t (Just a)
        Just _ -> retry
```

```
newtype TMVar a = TMVar (TVar (Maybe a))
```

# Question

- Threads block on an MVar are woken up in FIFO order
  - This is the fairness guarantee
- When multiple threads are blocked on a TVar, which should be woken up?
  - Consider: who can make progress? Example:

```
do x <- takeTVar v  
when (x != 42) retry
```

- All threads retrying on a variable are woken up

# Choice

- Choose an alternative action if the first transaction calls `retry`
  - If the first action returns a result, that is the result of the `orElse`
  - If the first action retries, the second action runs
  - If the second action retries, the whole action retries
  - Since the result of `orElse` is also an STM action, you can a ``orElse`` b ``orElse`` c ``orElse`` ...

`orElse :: STM a -> STM a -> STM a`

# Accounts, re-revisited

- Suppose we want to withdraw from a second account if the first has insufficient funds

```
withdraw2 :: Int -> Account -> Account -> STM ()  
withdraw2 amount account1 account2 =  
    withdraw amount account1  
    `orElse`  
    withdraw amount account2
```

# **STM as a building block (I)**

---

Asynchronous computations

# Asynchronous computations, revisited

- The goal:
  - Run computations asynchronously and wait for the results
  - Cancel and *race* running computations
  - Interface:

```
data Async a

async  :: IO a -> IO (Async a)
wait   :: Async a -> IO a
poll   :: Async a -> IO (Maybe a)
cancel :: Async a -> IO ()
race   :: Async a -> Async b -> IO (Either a b)
```

# async

- Perform an action *asynchronously* and later wait for the results

```
data Async a = Async ThreadId (TMVar a)

async :: IO a -> IO (Async a)
async action = do
    var <- newEmptyTMVarIO
    tid <- forkIO $ do
        res <- action
        atomically $ putTMVar var res
    return (Async tid var)
```

# wait

- Wait for the computation to complete

```
waitSTM :: Async a -> STM a
waitSTM (Async _ var) = readTMVar var
```

```
wait :: Async a -> IO a
wait a =
    atomically $ waitSTM a
```

```
race :: Async a -> Async b -> IO (Either a b)
race a b =
    atomically $
        fmap Left (waitSTM a)
        `orElse`
        fmap Right (waitSTM b)
```